
Bagging Deep Q-Networks

Jamis Johnson, Lance Legal, Angus Ding

Department of Computer Science

Columbia University

New York, NY

[jamis.johnson, lw12110, ad3180]@columbia.edu

Abstract

Recent advancements by DeepMind in leveraging convolutional neural networks to represent the state space in Q-learning have dramatically improved the performance of algorithmic game playing, outperforming nearly all competing algorithms across 49 Atari 2600 games. In this paper, we propose a straight forward, distributed method for increasing the efficacy of deep Q-network agents through the well-known ensemble algorithm, bootstrap aggregation. We have developed software to orchestrate the deployment, training, and testing of this method on a cluster of GPU optimized, Amazon Web Service EC-2 virtual machines. We provide results demonstrating the effectiveness of training time on performance as well as ensemble size to performance.

1 Overview

Q-networks have been shown to outperform 43 state-of-the-art agents across a diverse array of 49 Atari 2600 games, often by extreme margins.

1.1 Deep Q-networks

Q-learning is a long-standing, model-free reinforcement learning algorithm used to find an optimal action-selection policy[1]. Deep Q-networks are a novel approach to Q-learning in which deep convolutional neural networks are used to reduce high-dimensional raw input to a set of possible actions[2]. The convolutional neural network approximates the optimal action-value function

$$Q^*(s, a) = \max_{\pi} \mathbb{E} [r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots | s_t = s, a_t = a, \pi]$$

which is the maximum sum of discounted rewards at time t . It is quite common to approximate Q^* using a linear approximator, but in this case a nonlinear convolutional neural network is used such that, given θ weights, $Q(s, a; \theta) \approx Q^*(s, a)$. Spatial, convolutional neural networks have drastically increased the accuracy of image recognition tasks in recent years are therefore a natural fit as a state-space function approximator of raw image data[3,4,5,6,7].

Memory replay

DeepMind employed a biologically-inspired memory replay mechanism that was integral to the success of Q-networks. Experiences, $e_t = (s_t, a_t, r_t, s_{t+1})$, are stored at every time step t , where data set $D_t = \{e_1, \dots, e_t\}$. Q-learning updates are applied during learning on uniformly sampled experiences, with the loss function

$$L_i(\theta_i) = \mathbb{E}_{(s,a,r,s') \sim U(D)} \left[\left(r + \gamma \max_{a'} Q(s', a'; \theta_i^-) - Q(s, a; \theta) \right)^2 \right]$$

θ_i are the parameters to the Q-network and θ_i^- are the parameters used in the previous iteration to calculate the target.

Network architecture

A straightforward network architecture was chosen, presumably to prop up the generality of deep Q-networks. The input layer receives a 84 X 84 X 4 image produced by a prerocessing step ϕ which is in place to handle Atari 2600 artifacts such as flickering, a phenomenon that occurs because certain sprites are placed on odd or even numbered frames. The hidden layers are layer out as follows. First: convolves 32 filters, using an 8 X 8 filter, and a stride of 4; Second: 64 filters, a 4 X 4 filter, and a stride of 2; Third: 64 filters, a 3 X 3 filter, and a stride of 1. Each layer is followed by a rectified nonlinear layer, $\max(0, x)$. The last hidden layer is fully connected with 512 rectified units, and the final output layer is a fully-connected linear layer with an output per available action.

1.2 Bagging

Bootstrap aggregation is a standard technique for improving the stability and performance of statistical classifiers[8]. It is commonplace to employ this method when modeling using decision trees, but has also been shown to produce good results when applied to artificial neural networks[9,10].

Bagging is a process of averaging across an ensemble of classifiers $f^{*n}(x)$, where $n \in [1, N]$ and training data consists of N Bootstrap samples, such that

$$f_{bag}(x) = \frac{1}{N} \sum_{n=1}^N f^{*n}(x)$$

An intuitive way to think about bagging is: "the wisdom of crowds."

Bagging Q-networks

In a supervised learning environment you would train N neural networks, f_{*n} , on bootstrapped data sets, $Z_n = \{(x_1, y_1), \dots (x_b, y_b)\}_n$. But Q-networks do not explicitly operate within a supervised learning framework. Input data are dependent on the previous frame and the associated action taken. Because Q-networks are ϵ -greedy, random actions are injected regularly into the learning process. This allows us to train multiple networks, with the same initial parameters, in parallel, to achieve the same effect as bagging would in a supervised learning environment.

1.3 Alternate Approaches Considered

We initially pursued multiple avenues of optimization and tuning. Below summarizes those endeavors.

Multi-GPU Training Performance Speedup

Optimizing the deep Q-network code to utilize multiple GPUs on one machine was appealing because the code released by DeepMind was written using deprecated Torch libraries which had since been updated for multi-GPU processing. We decided to implement Krizhevsky's "weird trick" for optimizing the training of convolutional neural networks across machine-local GPUs: implementing data parallelism for the convolutional layers and model parallelism for the fully connected layers[11]. But this eventually proved to be unfruitful when we realized that completely saturating the memory and compute power of a single NVIDIA K520 GPU would not be straight forward. There is room however for work to be done in augmenting DeepMind's now dated Torch code to be concurrent by default.

Augmenting the Network Architecture

We toyed with the idea of replacing the straight forward network architecture with a more complex network architecture, some of which claim to have surpassed or matched human level image recognition[4][6][12]. Another option was parameter tuning, simply adjusting the depth, width, stride, and convolution size of the existing hidden layers to see what returned the best results. This idea was cast aside in favor of using the current initial network parameters and bagging the networks because we felt with the limited time allotted we had a higher probability of completion. With that said, we are extremely interested in applying these cutting edge architectures to deep Q-networks.

2 Implementation

2.1 Cluster Architecture

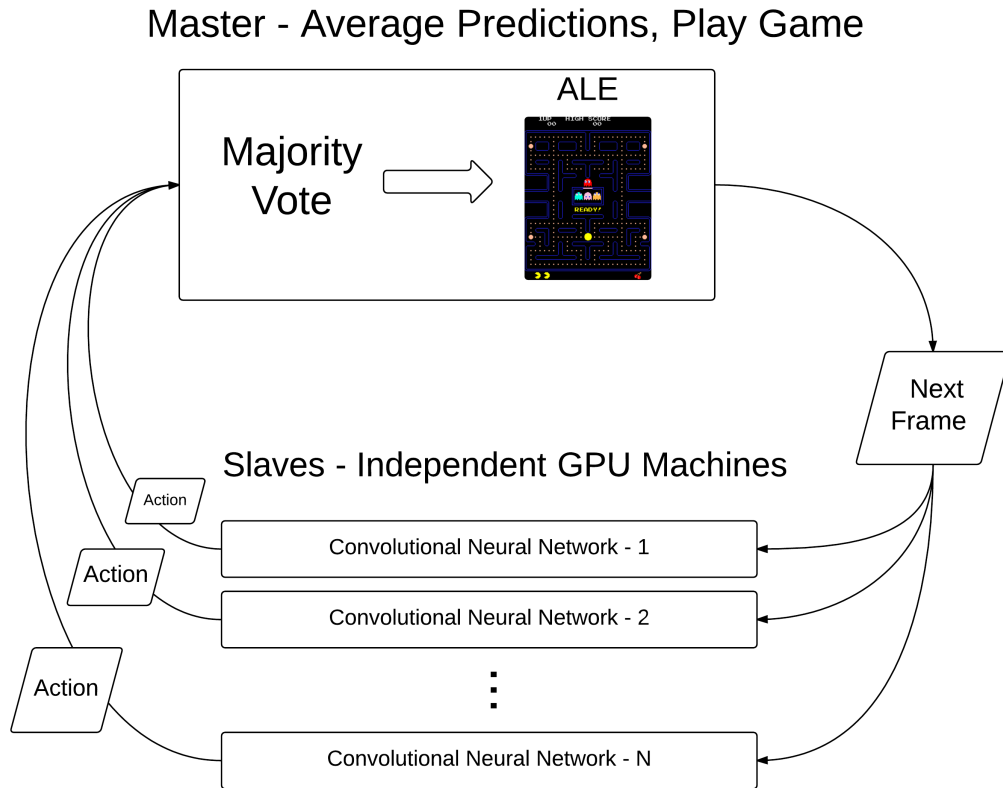


Figure 1: Cluster Architecture

Table 1: TODO: Title

Config	Space	Breakout	Pacman
3x6	6	5	5
3x10	6	5	5
3x18	6	5	5

3 Results




Figure 2: Effect of Ensemble Size on Performance.




Figure 3: Effect of Training Time on Ensembles.

4 Conclusion

Acknowledgments

Thanks Mom

References

- [1] Watkins, Christopher JCH, and Peter Dayan. "Q-learning." *Machine learning* 8.3-4 (1992): 279-292.
- [2] Mnih, Volodymyr, et al. "Human-level control through deep reinforcement learning." *Nature* 518.7540 (2015): 529-533.
- [3] Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E. Hinton. "Imagenet classification with deep convolutional neural networks." *Advances in neural information processing systems*. 2012.
- [4] Szegedy, Christian, et al. "Going deeper with convolutions." *arXiv preprint arXiv:1409.4842* (2014).
- [5] Simonyan, Karen, and Andrew Zisserman. "Very deep convolutional networks for large-scale image recognition." *arXiv preprint arXiv:1409.1556* (2014).
- [6] He, Kaiming, et al. "Delving deep into rectifiers: Surpassing human-level performance on imagenet classification." *arXiv preprint arXiv:1502.01852* (2015).
- [7] Sermanet, Pierre, et al. "Overfeat: Integrated recognition, localization and detection using convolutional networks." *arXiv preprint arXiv:1312.6229* (2013).
- [8] Breiman, Leo. "Bagging predictors." *Machine learning* 24.2 (1996): 123-140.
- [9] Schwenk, Holger, and Yoshua Bengio. "Boosting neural networks." *Neural Computation* 12.8 (2000): 1869-1887.

- [10] Ha, Kyoungnam, Sungzoon Cho, and Douglas MacLachlan. "Response models based on bagging neural networks." *Journal of Interactive Marketing* 19.1 (2005): 17-30.
- [11] Krizhevsky, Alex. "One weird trick for parallelizing convolutional neural networks." *arXiv preprint arXiv:1404.5997* (2014).
- [12] Ioffe, Sergey, and Christian Szegedy. "Batch normalization: Accelerating deep network training by reducing internal covariate shift." *arXiv preprint arXiv:1502.03167* (2015).
- [] Dean, Jeffrey, et al. "Large scale distributed deep networks." *Advances in Neural Information Processing Systems*. 2012.
- [] Collobert, Ronan, Koray Kavukcuoglu, and Clment Farabet. "Torch7: A matlab-like environment for machine learning." *BigLearn, NIPS Workshop*. No. EPFL-CONF-192376. 2011.